


**「aws serverless apps」**

# aws serverless apps

## introduction



「 what is  
serverless? 」

# parts of serverless apps



## JavaScript

All of the dynamic behavior is provided by JavaScript running in the user's browser. The JavaScript itself is static, though.



## API

Data is written to and retrieved from an API hosted by a cloud provider. In our case, that provider is AWS API Gateway.



## Markup

Finally, the HTML markup ties the whole thing together. The markup is static and cacheable, so the initial page load for your users is fast.

# how do they work together?

a user comes to your site and loads static **HTML, CSS, and JavaScript**

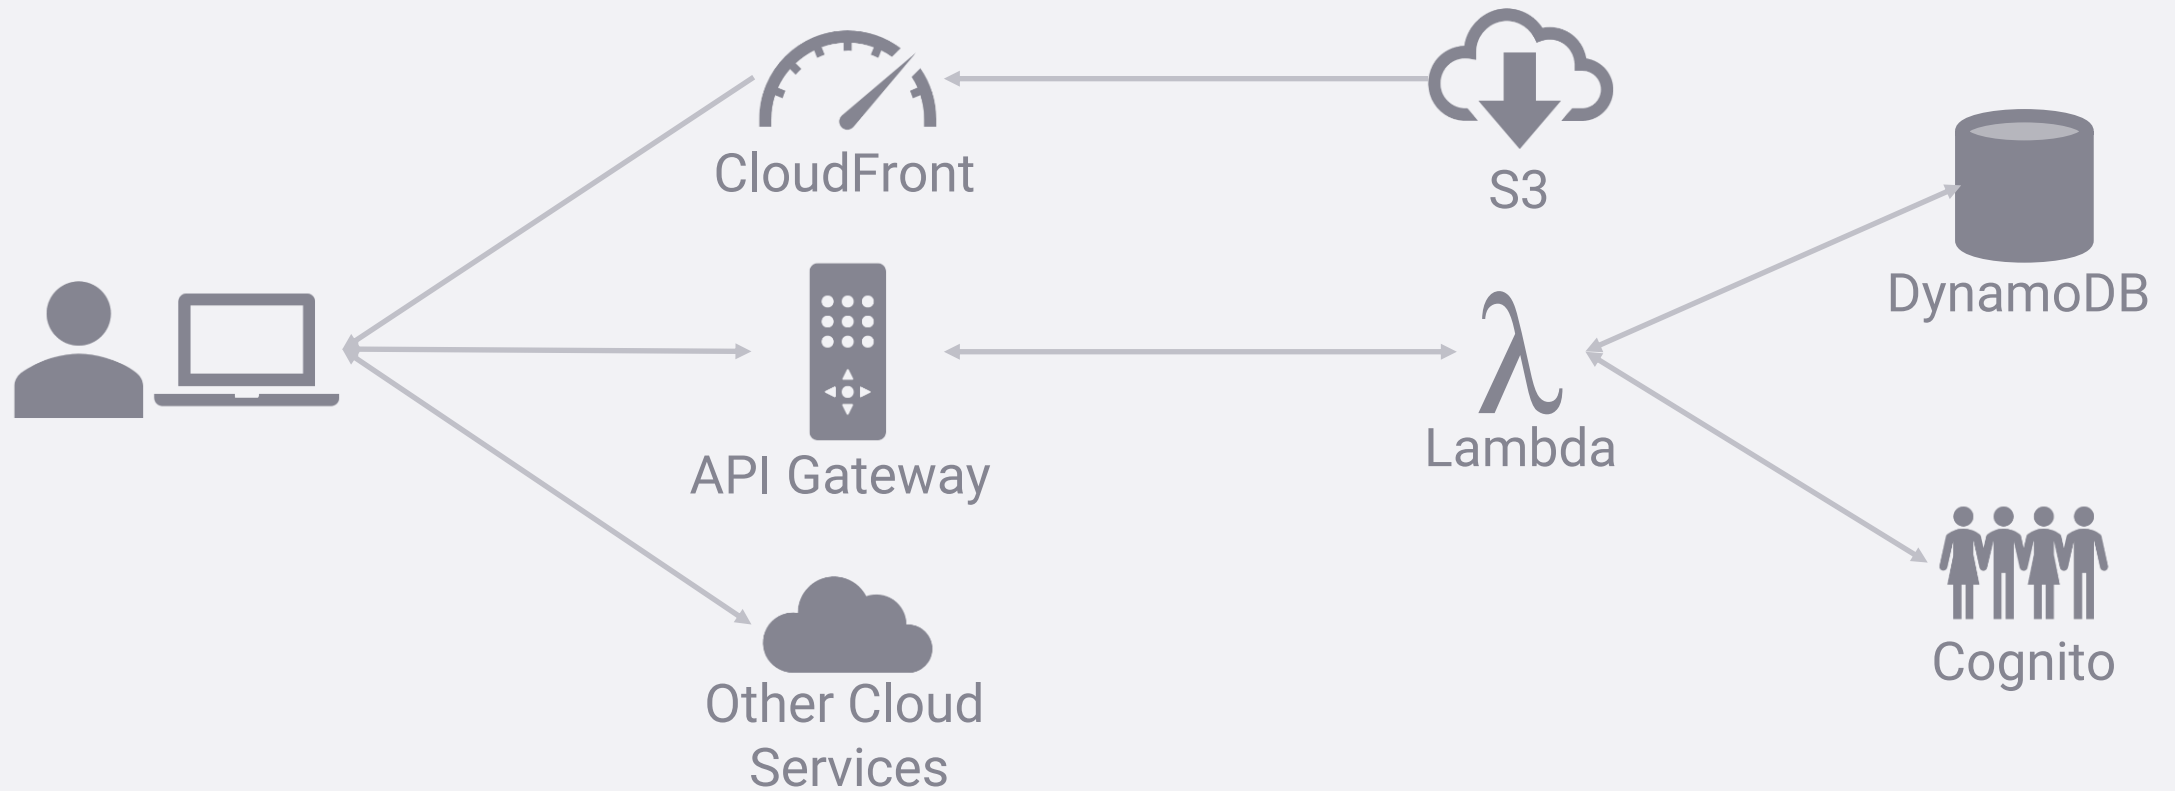
static parts of the page are in **HTML** while dynamic parts of the page are added by **JavaScript**

**JavaScript** retrieve dynamic data from the **API**

**JavaScript** writes data to the **API**

**JavaScript** handles some aspects of navigation

maybe a picture?



「 why go  
serverless? 」

# why not just run a server?

server management (patches, monitoring, hardware failures, etc.)

servers can be cheap, but scaling gets expensive really fast

you don't pay for processing time you don't use

easier to split up development between front-end and back-end



**aren't s3, lambda, etc.  
servers?**

yes, but...

they're not yours

you're not the only one on them

amazon has people with pagers to keep them working

# good and bad use cases

good:

dynamic applications with lots of user interaction  
most of the content in the UI is specific to the user

bad:

large applications with lots of data that would have to be loaded  
anything where the time to load markup, JS, and data would create a very slow UI

demo  
app



「getting started」

# getting started



# sign up for AWS

Sign up at: <https://portal.aws.amazon.com/billing/signup>

Free tier information: <https://aws.amazon.com/free/>

Services we will use:

- S3
- API Gateway
- Lambda
- DynamoDB
- Cognito
- CloudFront
- IAM
- CloudWatch





「understanding S3」

# S3 basics

S3 = Simple Storage Service

Key-Blob store

Eventually consistent

Extremely durable

**S3 is not...**

a file system

# one more thing...

AWS policies define access control for:

- users
- other AWS services
- resources

# make life easier

```
{
  "Version": "2012-10-17",
  "Id": "Policy1497053408897",
  "Statement": [
    {
      "Sid": "Stmt1497053406813",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::<your bucket name>/*"
    }
  ]
}
```



# 「JavaScript syntax」

# JavaScript syntax

JavaScript is a C-style language in most respects:

- control structures
- variable naming rules
- accessing member values of an object (e.g. foo.bar)

JavaScript is not always C-like:

- variables have no type
- functions are objects
- ===



# JavaScript syntax oddities

== - equivalent values

=== - equal values of the same type

Example:

1 == "1" is true

1 === "1" is false

# JavaScript syntax oddities

null vs. undefined

null – an object of type null

undefined – a value of type undefined

# JavaScript declarations

```
var foo = 'bar';  
var baz = "bar";
```

```
var object = {  
  p1: 1,  
  p2: 2,  
  'this needs to be quoted': 'Hello, world!'  
};
```

object.p1 is the same as object['p1']

# JQuery

\$ is JQuery

\$.get()

\$.post()

\$('.class') – CSS class selector

\$('#id') – ID selector

\$('div') – tag selector

# Useful references

JavaScript: <https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference>

JQuery: <http://api.jquery.com/>



# 「JavaScript parallelism」

# classic parallelism



Scenario:

You need to make two independent HTTP requests as fast as possible

Solution:

Create two threads to make the requests  
Use some sort of synchronization to determine when they're done  
Keep doing what you need to do



# is that bad?

No, it's perfectly normal!

Good:

Assuming you have enough processors and memory, you can be as efficient as possible.  
The threads are not lies. They're real, and you can make them do what you want.  
It's very powerful.

Bad:

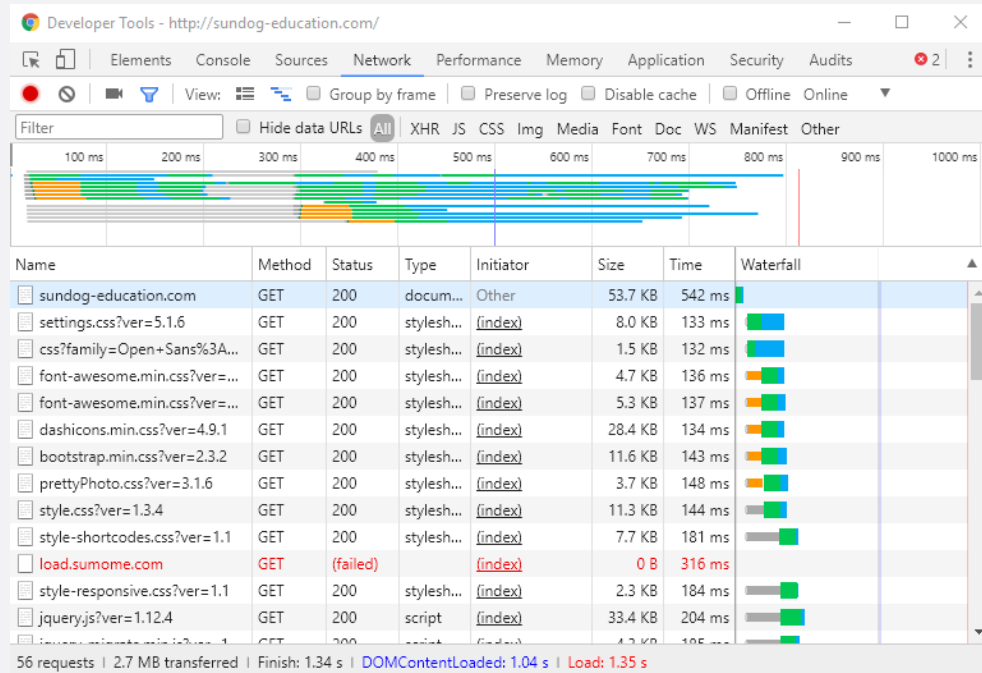
It's very complicated.  
Most people do synchronization poorly.  
There are entire books on the subject of proper multithreading.

# JavaScript threading

JavaScript was made to run in the browser.  
The DOM can only be manipulated in a single thread.

Since JavaScript was originally intended to manipulate the DOM, there's no reason to introduce multithreading.

# what about AJAX?



Exactly!

Browsers already do a bunch of HTTP requests in parallel!

# JavaScript is single-threaded!

...ish?

All your code runs on a single thread.

Browsers and Node.js manage other threads  
that do things for you.



# callbacks

## jQuery.get()

Categories: [Ajax](#) > [Shorthand Methods](#)

**jQuery.get( url [, data ] , success ) [, dataType ]**

**Description:** Load data from the server using an asynchronous HTTP GET request.

**url**  
Type: [String](#)  
A string containing the URL to which the request is sent.

Gives a description of the Time to Live (TTL) status on the specified table.

`getItem(params = {} , callback) ⇒ AWS.Request`  
The GetItem operation returns a set of attributes for the item with the given primary key.

`listBackups(params = {} , callback) ⇒ AWS.Request`  
List backups associated with an AWS account.

`listGlobalTables(params = {} , callback) ⇒ AWS.Request`  
Lists all the global tables.

`listTables(params = {} , callback) ⇒ AWS.Request`  
Returns an array of table names associated with the current account and endpoint.

`listTagsOfResource(params = {} , callback) ⇒ AWS.Request`  
List all tags on an Amazon DynamoDB resource.

`putItem(params = {} , callback) ⇒ AWS.Request`  
Creates a new item, or replaces an old item with a new item.

# effects of event loop

Dependent calls require a lot of functions

Bookkeeping is everything for parallel behaviors

No synchronization

「that's it!」





# exercise #1

- a) Add a message from “Student” to Frank
- b) Add a new conversation



**「exercise solution」**

# exercise #1

## solution

```
2.json:
{
  "id": "2",
  "participants": ["Student", "Frank"],
  "messages": [
    {
      "sender": "Frank",
      "time": 1512246342159,
      "message": "This is Frank. I'm also sending you a message."
    },
    {
      "sender": "Student",
      "time": 1512247342159,
      "message": "Hello, Frank!"
    }
  ]
}
```

# exercise #1

## solution

```
conversations.json:
[
  {
    "participants": ["Student", "Brian"],
    "id": "1"
  },
  {
    "participants": ["Student", "Frank"],
    "id": "2"
  },
  {
    "participants": ["Student", "Alice"],
    "id": "hello-world"
  }
]
```

# exercise #1

## solution

```
hello-world.json
{
  "id": "hello-world",
  "participants": ["Student", "Alice"],
  "messages": [
    {
      "sender": "Alice",
      "time": 1512246299194,
      "message": "I'm Alice. I usually talk to Bob."
    }
  ]
}
```







# what is lambda?

Function running on AWS

Options:

- Java
- Python
- JavaScript (Node.js)

# what makes lambda serverless?

Feature	Lambda	EC2
Time to spin up	milliseconds	seconds to minutes
Billing increments	100ms	1s
Configuration	function	operating system
Scaling unit	parallel function executions	instances
Maintenance	AWS	AWS and you



# 「IAM and policies」

# IAM? i am what?

IAM – Identity and Access Management

users/groups

roles

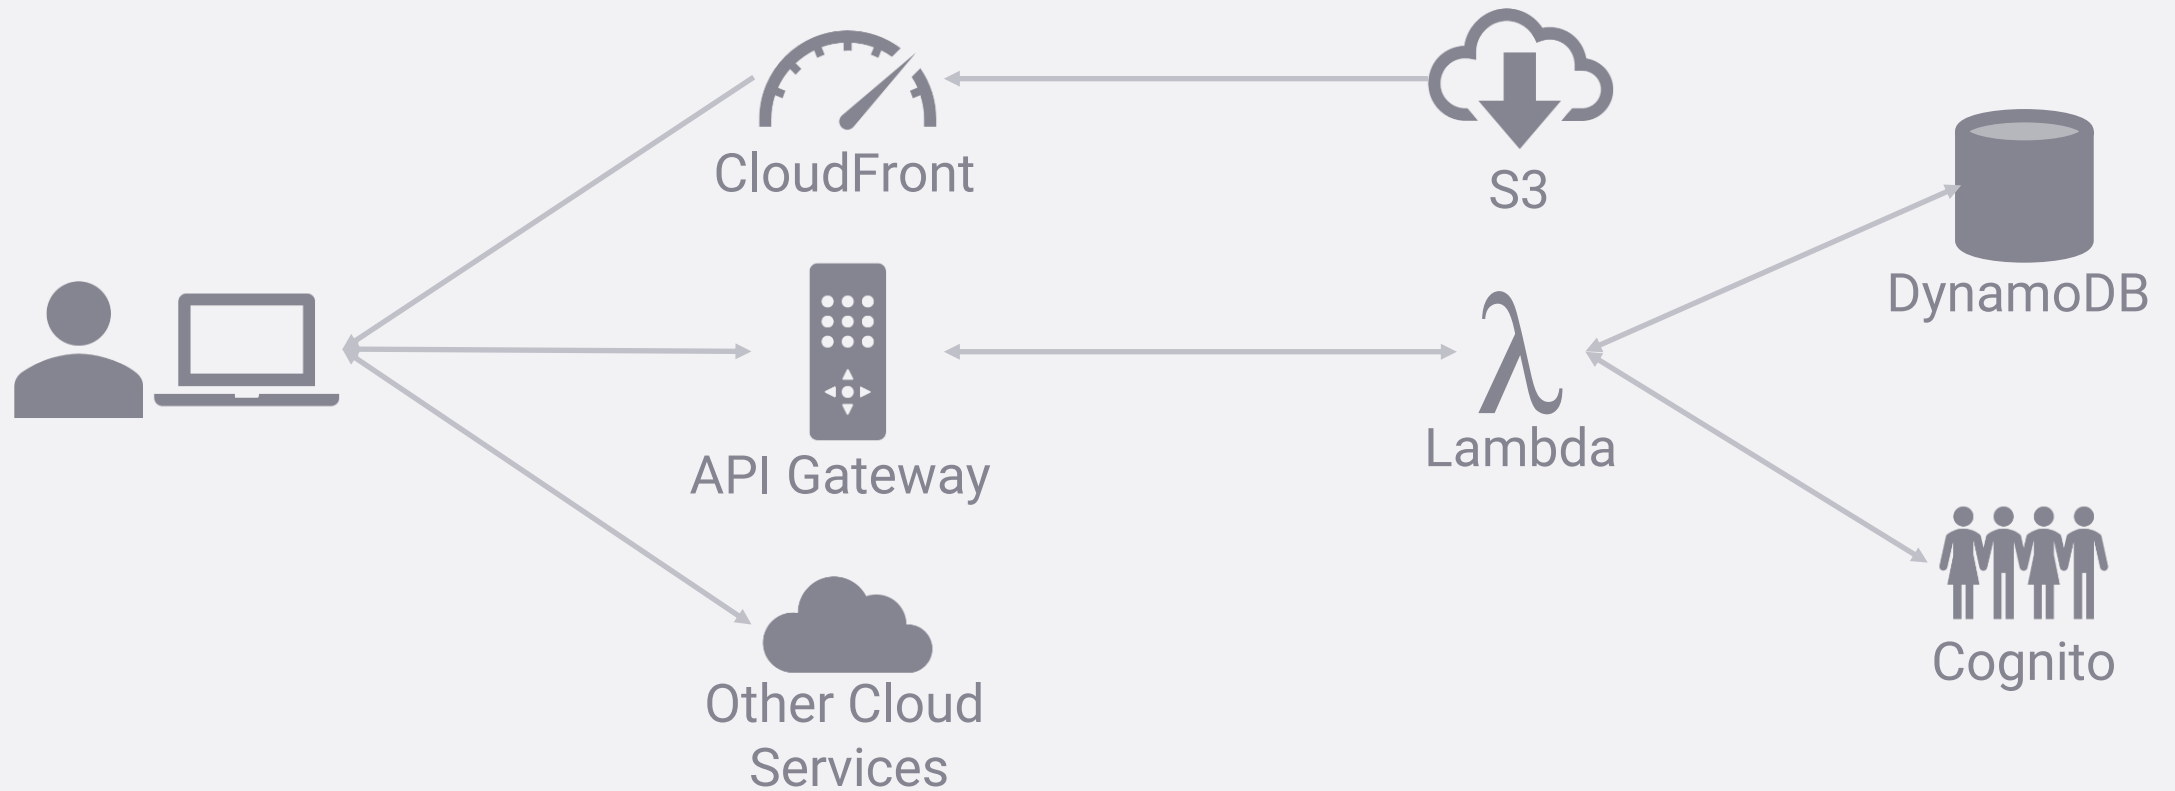
policies

identity providers

user account settings

encryption keys

# where do we need roles?





「lambda triggers」



# so many options

API Gateway  
AWS IoT  
Alexa Skills  
Alexa Smart Home  
CloudWatch Events  
CloudWatch Logs  
CodeCommit  
Cognito Sync  
DynamoDB  
Kinesis  
S3  
SNS

| but wait...there's  
more!


Lambda API  
CloudFront

# api gateway

proxy



mapping





「cors」

**cors?**

cross-origin resource sharing

# long ago, on the internet

load this image from bar.com in an <img> tag?

sure

website  
foo.com



load this JS from baz.com in a <script> tag?



browser

sure

make an XMLHttpRequest to baz.com?

NO!!!

# someone got clever

load this JS from baz.com in a `<script>` tag?

website  
foo.com



here's a new `<script>` tag to load from baz.com

sure



browser

sure



# new standard: cors

load this JS from baz.com in a <script> tag?

website  
foo.com



make an XMLHttpRequest to baz.com?

sure



browser

maybe?

maybe?

OPTIONS /api  
Origin  
Access-Control-Request-Method  
Access-Control-Request-Headers

browser




Access-Control-Allow-Origin  
Access-Control-Allow-Methods  
Access-Control-Allow-Headers  
Access-Control-Max-Age  
[Access-Control-Allow-Credentials]



baz.com

# then what?

browser  POST /api  
Origin

Access-Control-Allow-Origin  baz.com



「dynamodb」

# nosql

“non SQL” or “Not only SQL”

basically a key-value storage system

typically not ACID

extremely scalable

# dynamodb tables



# types of attributes

name	type
B and BS	binary (and set)
BOOL	boolean
L	list
M	map
N and NS	number (and set)
NULL	null
S and SS	string (and set)



# retrieving items

get item

full key



item

query

partial key



results

scan

filter



results

## a note on pagination



# provisioned throughput

1 read unit = 1 consistent read up to 4KB per second

1 read unit = 2 eventually consistent reads up to 4KB per second

1 write unit = 1 write up to 1KB per second

# that's the basics

table structures  
data types  
how to get items  
pagination  
provisioned capacity



「dynamodb vs. s3」

# blobs vs. values

001010101001010  
101010101010101  
010101010101010  
010101010101010  
000000011010101

key	value
key 1	stuff
key 2	other stuff

# limits

limit	S3	DynamoDB
record size	5TB	400KB
request rate	basically unlimited	80k capacity units in Virginia, 20k everywhere else
total size	unlimited	unlimited



**「consistency!」**



## a few things to try

1. Make the Lambda function reject unsupported HTTP methods
  - A. Create some new test configurations for these cases
2. Remove S3 access for your Lambda function



「λ proxy pain」

# wait...what?

poorly-documented event object

having to worry about the details of HTTP requests

input comes from multiple places

over 200 lines of JavaScript to manage

```
index.js
1 'use strict';
2
3 var AWS = require('aws-sdk');
4
5 var dynamo = new AWS.DynamoDB();
6
7 exports.handler = function (event, context, callback) {
8
9     const done = function (err, res) {
10         callback(null, {
11             statusCode: err ? '400' : '200',
12             body: err ? JSON.stringify(err) : JSON.stringify(res),
13             headers: {
14                 'Content-Type': 'application/json',
15                 'Access-Control-Allow-Origin': '*'
16             }
17         });
18     };
19
20     var path = event.pathParameters.proxy;
21
22     if (path === 'conversations' && event.httpMethod === 'GET') {
23         dynamo.query({
24             TableName: 'Chat-Conversations',
25             IndexName: 'Username-ConversationId-index',
26             Select: 'ALL_PROTECTED_ATTRIBUTES',
27             KeyConditionExpression: 'Username = :username',
28             ExpressionAttributeValues: {':username': {s: 'Student'}}
29         }, function (err, data) {
30             handleIdQuery(err, data, done, [], 'Student');
31         });
32     } else if (path.startsWith('conversations/')) {
33         var id = path.substring('conversations/'.length);
34         switch(event.httpMethod) {
35             case 'GET':
36                 dynamo.query({
37                     TableName: 'Chat-Messages',
38                     ProjectionExpression: '#1, Sender, Message',
39                     ExpressionAttributeNames: {'#1': 'Timestamp'},
40                     KeyConditionExpression: 'ConversationId = :id',
41                     ExpressionAttributeValues: {'id': {s: id}}
42                 }, function (err, data) {
43                     loadMessages(err, data, id, [], done);
44                 });
45                 break;
46             case 'POST':
47                 dynamo.putItem({
48                     TableName: 'Chat-Messages',
49                     Item: {
50                         ConversationId: {s: id},
51                         Timestamp: {
52                             N: "" + new Date().getTime()
53                         },
54                         Message: {s: event.body},
55                         Sender: {s: 'Student'}
56                     }
57                 }, done);
58                 break;
59             default:
60                 done('No cases hit');
61                 break;
62         }
63     } else {
64         done('No cases hit');
65     }
66 };
67
68 function loadMessages(err, data, id, messages, callback) {
69     if (err === null) {
70         data.Items.forEach(function (message) {
71             messages.push({
72                 sender: message.Sender.S,
73                 time: Number(message.Timestamp.N),
74                 message: message.Message.S
75             });
76         });
77         if (data.LastEvaluatedKey) {
78             dynamo.query({
79                 TableName: 'Chat-Messages',
80                 ProjectionExpression: '#1, Sender, Message',
81                 KeyConditionExpression: 'ConversationId = :id',
82                 ExpressionAttributeNames: {'#1': 'Timestamp'},
83                 ExpressionAttributeValues: {'id': {s: id}},
84                 ExclusiveStartKey: data.LastEvaluatedKey
85             }, function (err, data) {
86                 loadMessages(err, data, id, messages, callback);
87             });
88         } else {
89             loadConversationDetail(id, messages, callback);
90         }
91     }
92 }
```

# is there anything good about it?

it's easy

Lambda startup cost can be lower

less management

Duration	1336.80 ms
----------	------------

Duration	542.83 ms
----------	-----------

Duration	204.64 ms
----------	-----------

| basically...





# the power of api gateway

each resource-method combination can be its own function with its own lifecycle

each function can have simple, tailored input

ensuring API compatibility can occur without code

# the plan

resources and methods

api models

request and response flows



**「resources and methods」**

# http basics

```
GET / HTTP/1.1
User-Agent: curl/7.38.0
Host: www.google.com
Accept: */*
```

```
HTTP/1.1 200 OK
Date: Thu, 11 Jan 2018 03:46:43 GMT
Expires: -1
Cache-Control: private, max-age=0
Content-Type: text/html; charset=ISO-8859-1
P3P: CP="This is not a P3P policy! See
g.co/p3phelp for more info."
Server: gws
X-XSS-Protection: 1; mode=block
X-Frame-Options: SAMEORIGIN
Set-Cookie: ...
Accept-Ranges: none
Vary: Accept-Encoding
Transfer-Encoding: chunked
```

# methods



## Read

GET  
HEAD



## Write

POST  
PUT  
PATCH



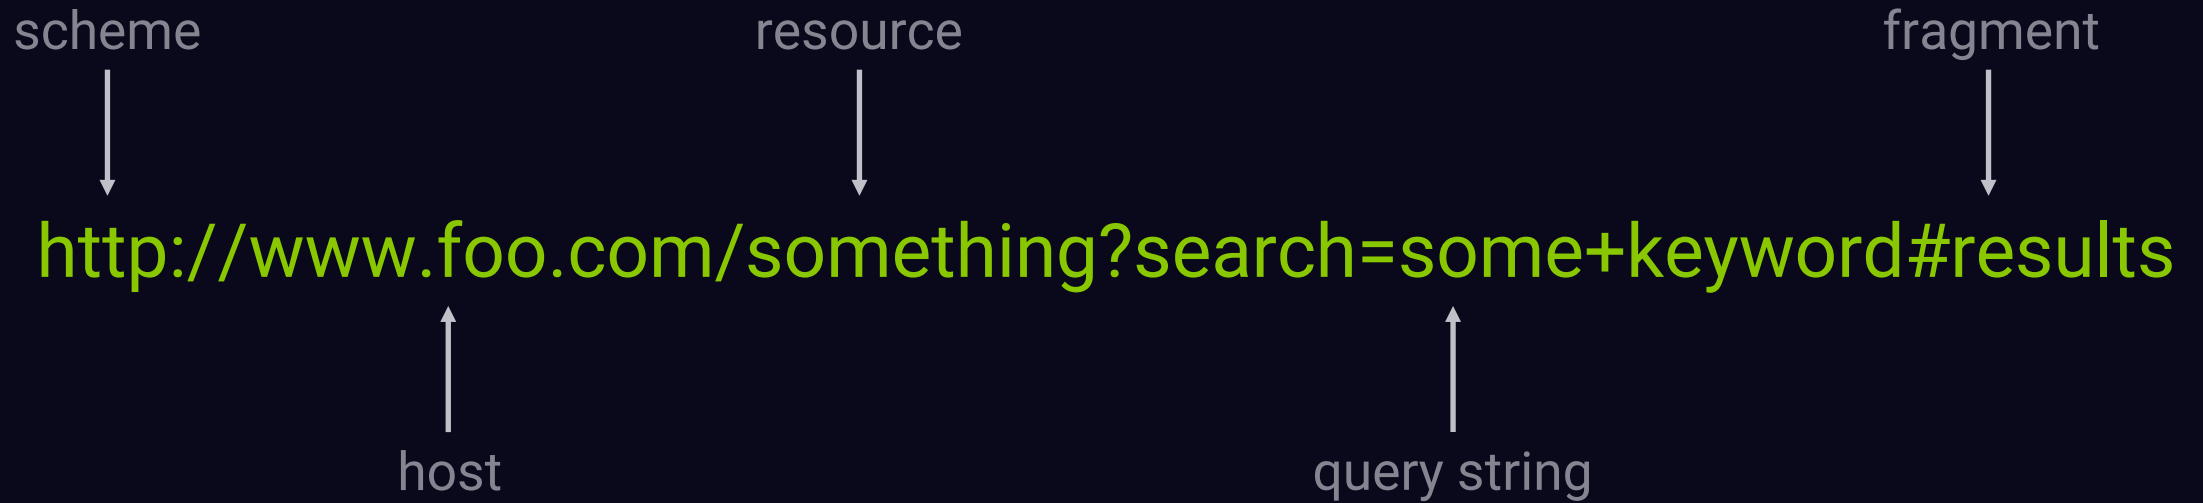
## Other

DELETE  
OPTIONS

# status codes

family	meaning
2xx	success
3xx	redirect
4xx	client error
5xx	server error

# parts of a URL



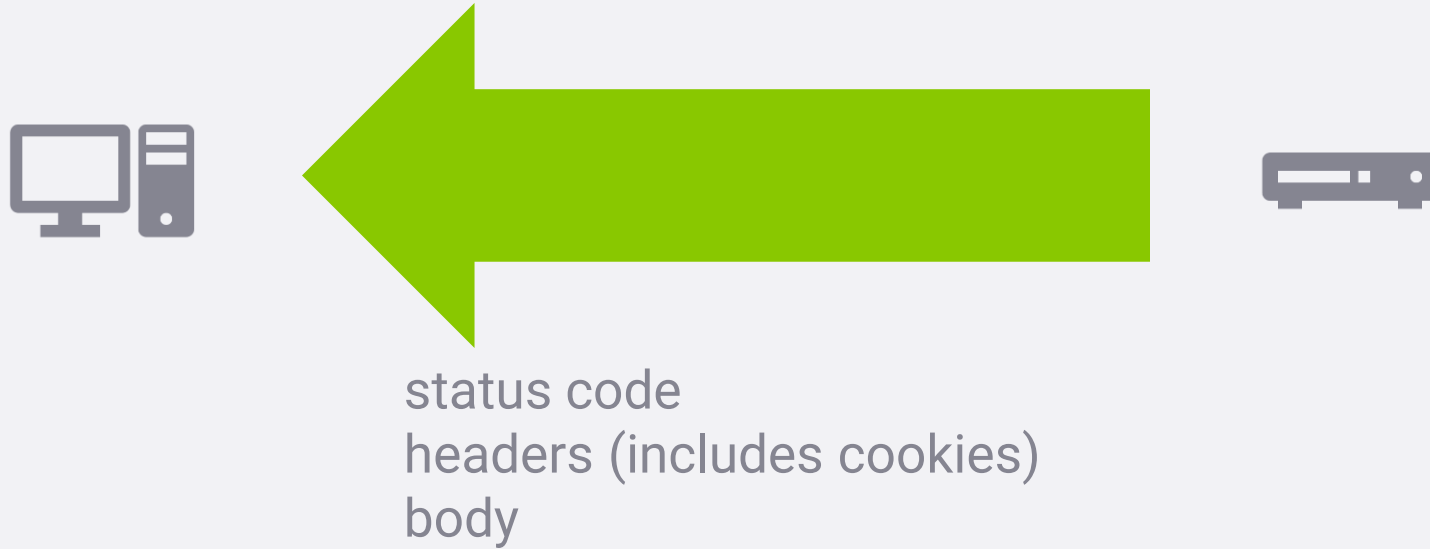


# request data



resource  
method  
query string  
headers (includes cookies)  
body (POST, PUT, PATCH)

# response data



# what goes where?

what?	request	response
resource	what data or behavior is required	N/A
query string	filtering and subset information	N/A
headers	processing directives	processing directives
body	input data	output data

# resources



## Fully-defined

Some resources are one of a kind even though they may change based on user, time, etc.

Example: /conversations

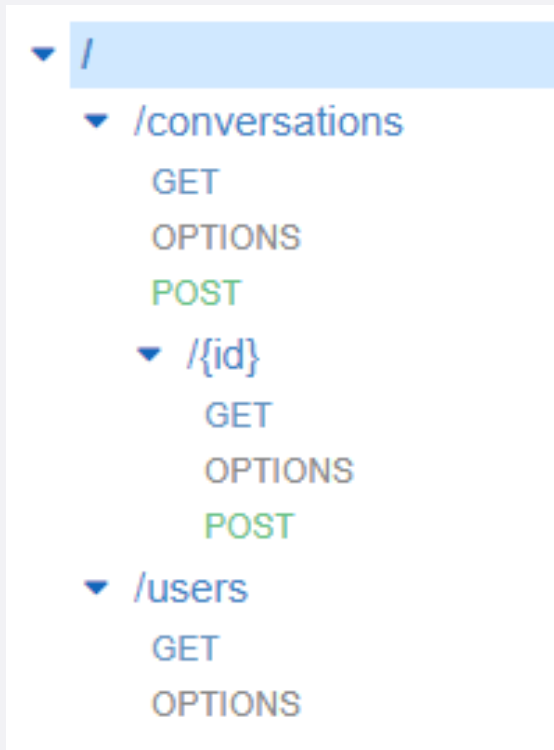


## Parameterized

Other resources contain data in the URL. There can be some large number of resources that follow a pattern.

Example: /conversations/{id}

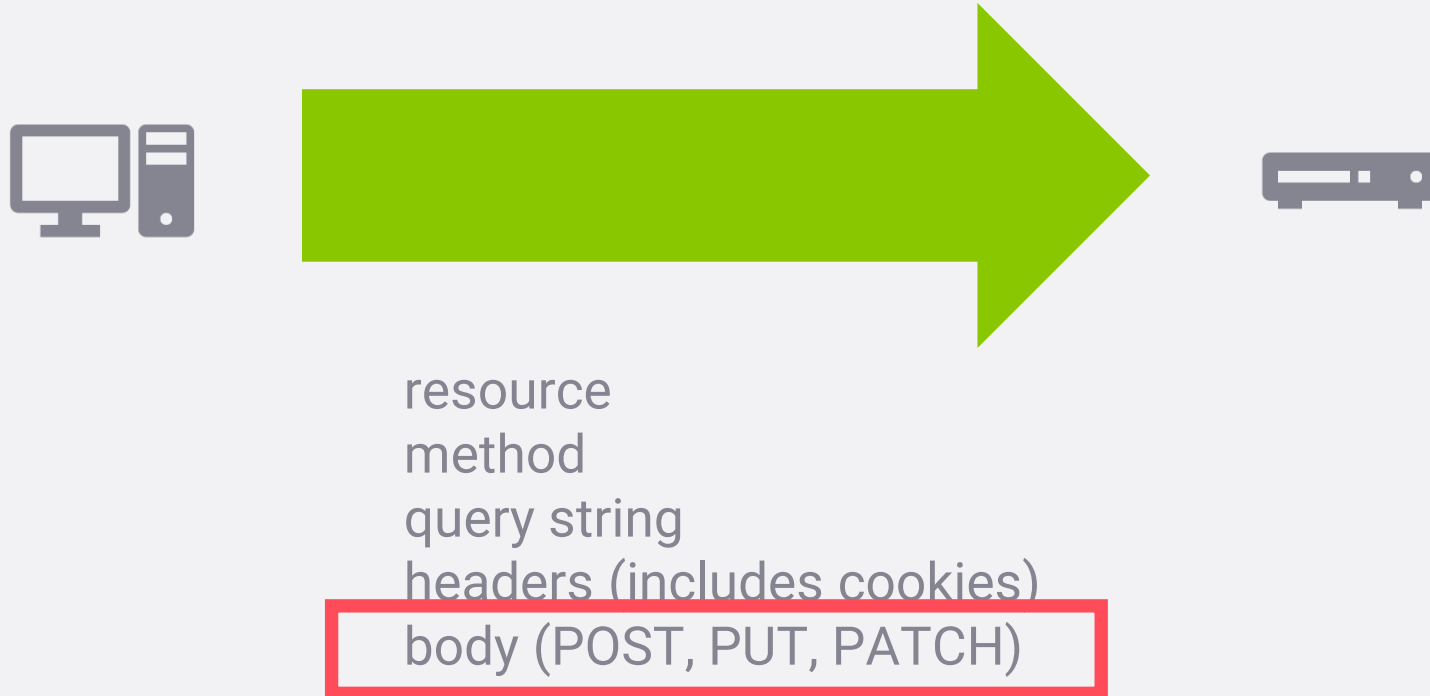
# resources have methods





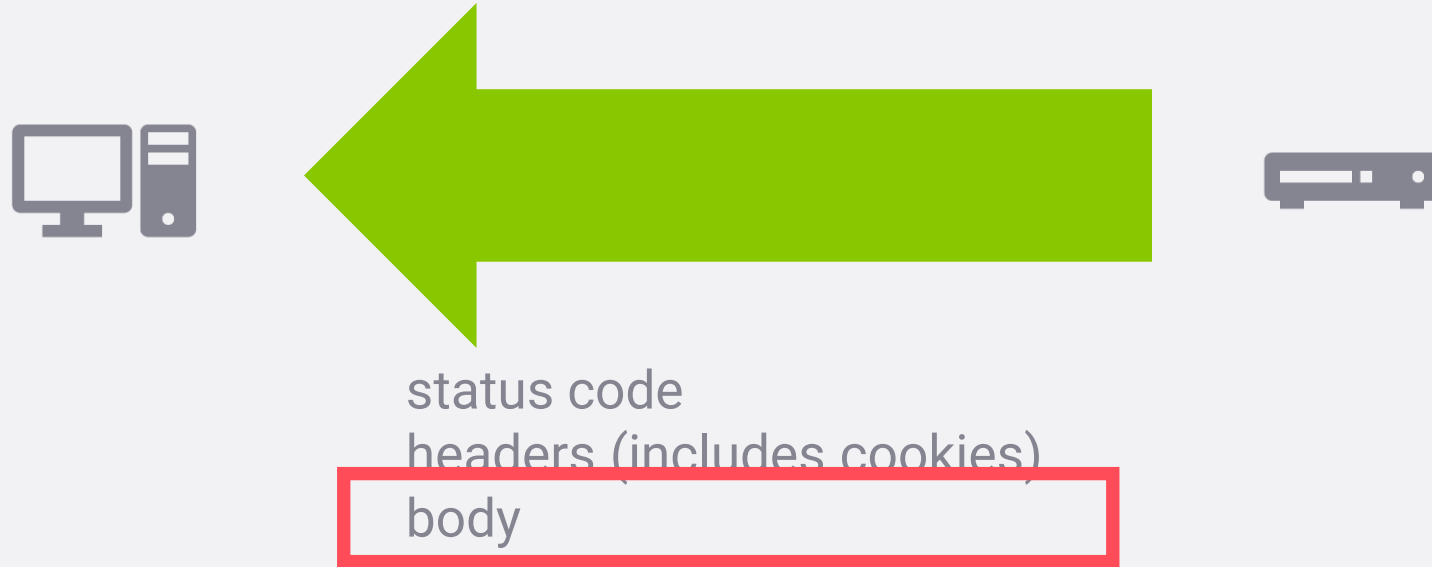
「models」

# request data





# response data



# models in api gateway

ChatAPI

Resources

Stages

Authorizers

Gateway Responses

Models

Documentation

Dashboard

Settings

# defining models

JSON schema

each model in its own schema document

# json schema primitive types

type	notes
null	
boolean	
object	properties defined by “properties”
array	elements defined by “items”
number	
string	

# examples

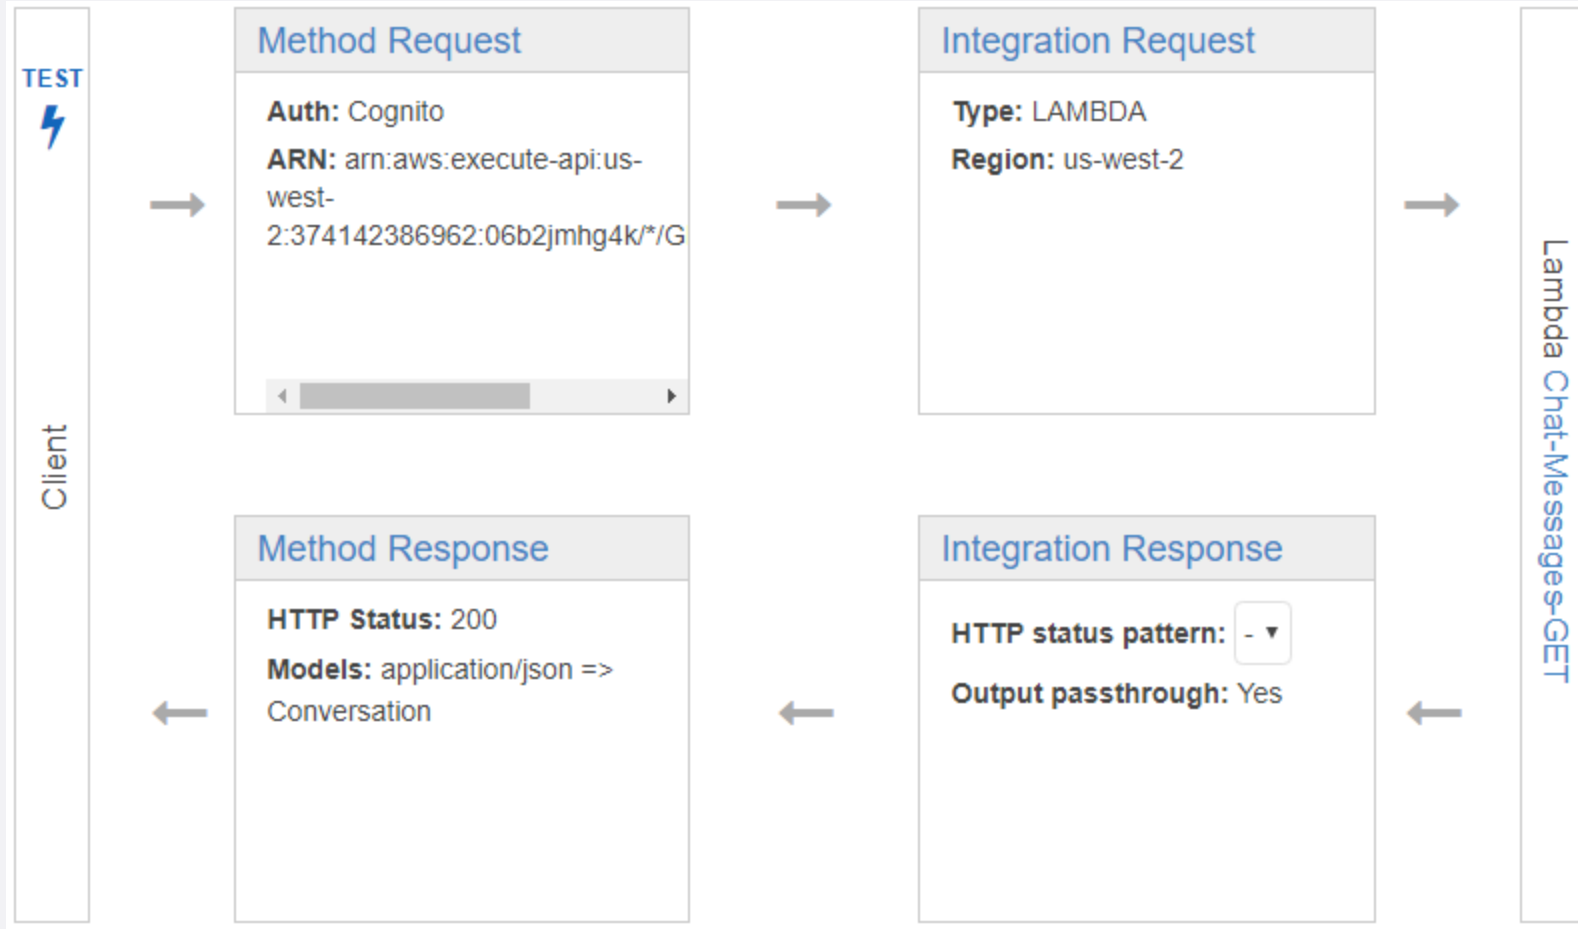
```
{  
  "type": "string"  
}
```

```
{  
  "type": "array",  
  "items": {  
    "type": "object",  
    "properties": {  
      "id": {  
        "type": "string"  
      },  
      "participants": {  
        "type": "array",  
        "items": {  
          "type": "string"  
        }  
      },  
      "last": {  
        "type": "number",  
        "format": "utc-millsec"  
      }  
    }  
  }  
}
```



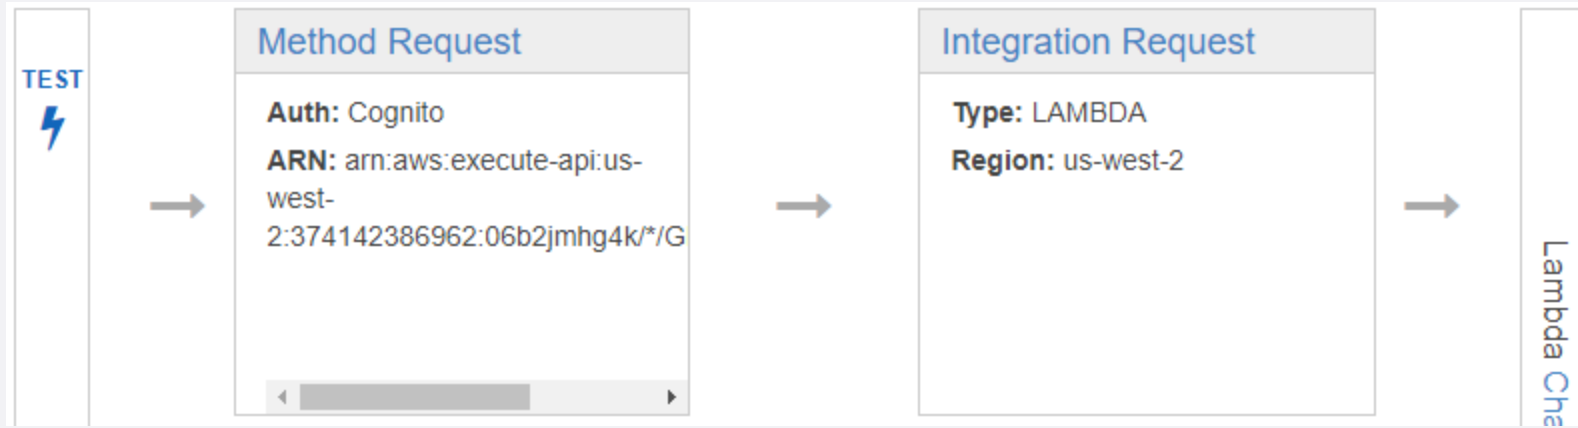
「request flow」

# overall flow





# request flow



authorization  
validation  
caching  
input:

- request path variables
- query string parameters
- request headers
- request body

integration type  
remapping of input

# authorization

AWS IAM  
API key  
Cognito  
Custom

# integration

**Integration type**

- ☒ Lambda Function ⓘ
- ☐ HTTP ⓘ
- ☐ Mock ⓘ
- ☐ AWS Service ⓘ
- ☐ VPC Link ⓘ

# lambda and mock



## Lambda

Supports proxy and non-proxy  
modes

Just calls a Lambda function



## Mock

Always returns the same thing

# http, aws, and vpc

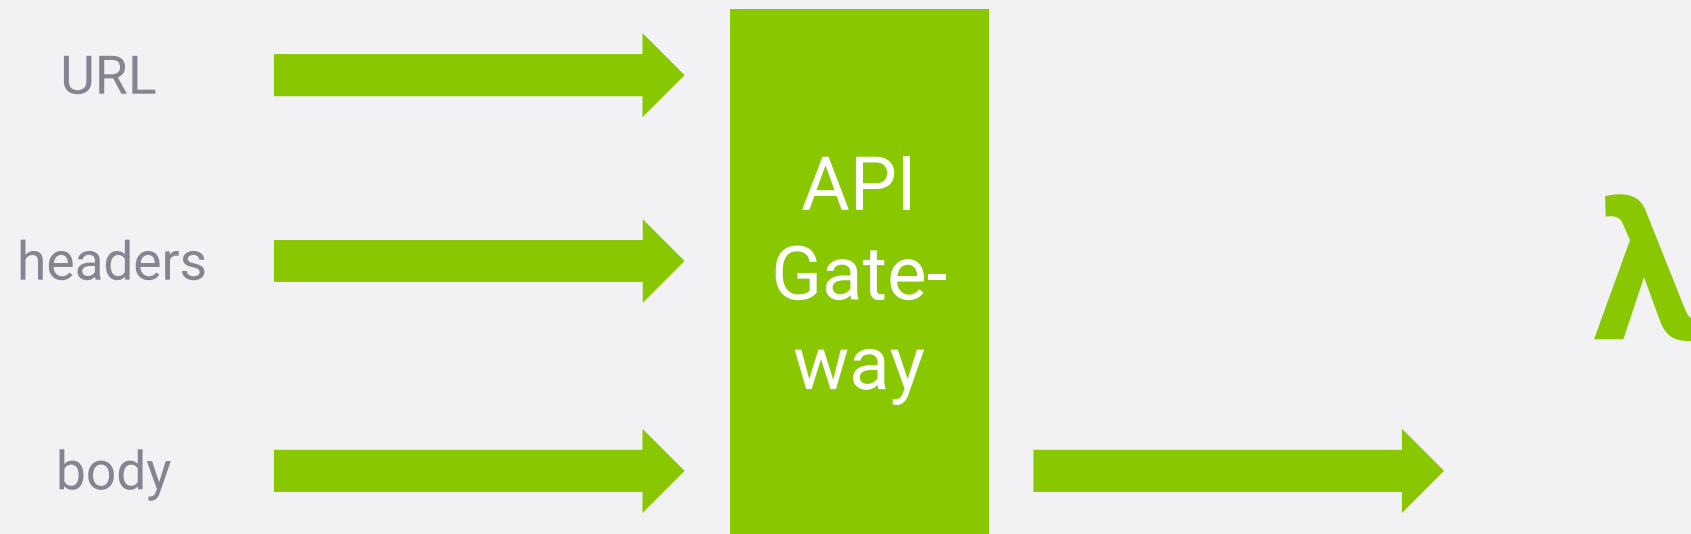
all very similar

default is effectively passthrough

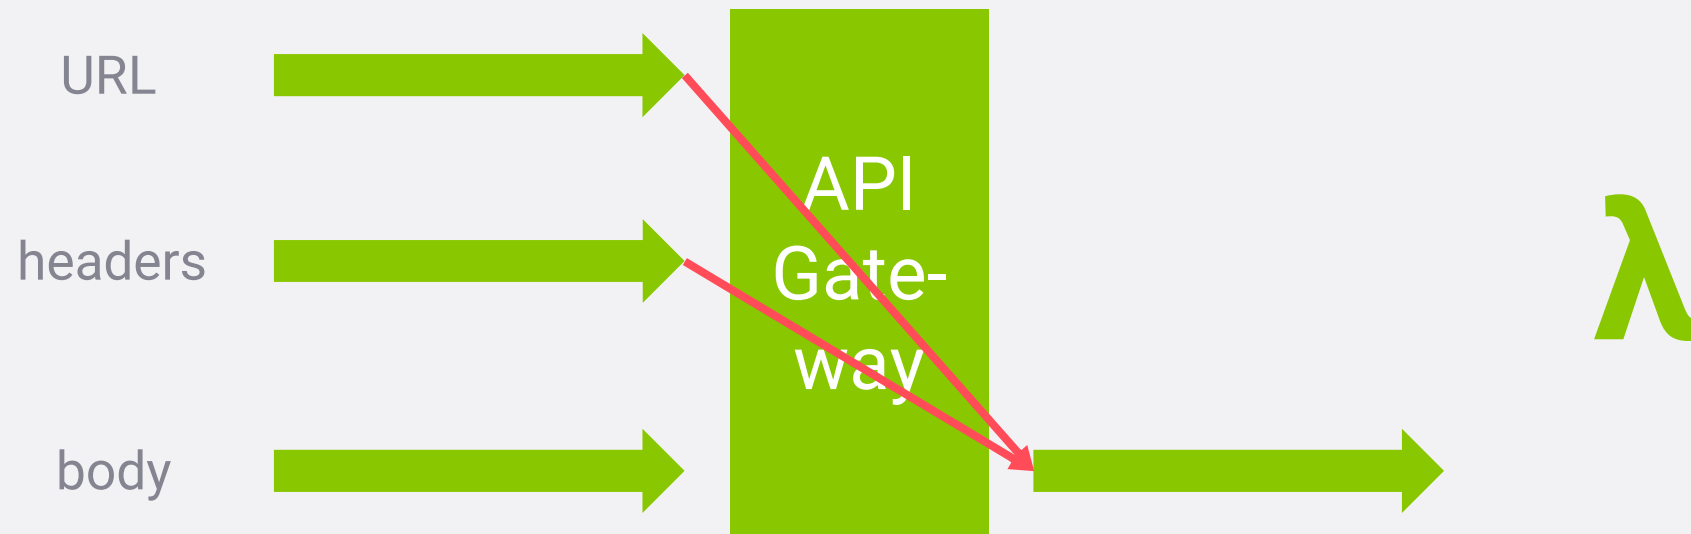
remapping can be done on almost anything

also, proxy mode

# input mapping




# input mapping



「response flow」



# overall flow



# response flow



response definitions

response mapping  
status code mapping



「stages」

# what do they do?

snapshots of the API

public URL

operational settings

generated clients

generated swagger

Canary

# operational stuff



## Caching

CloudFront-based caching



## Throttling

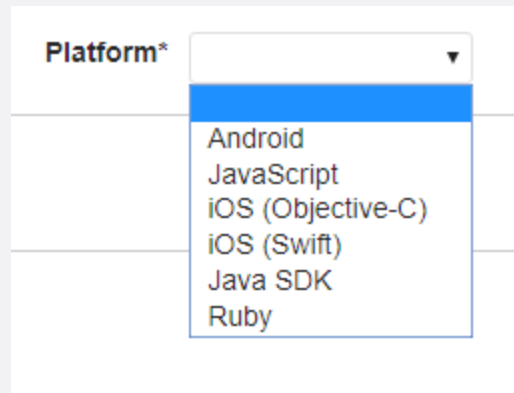
Request rate with bursting



## Variables

Define variables for the stage

# sdk generation and swagger



Export as Swagger




Export as Swagger + API  
Gateway Extensions




Export as Swagger + Postman  
Extensions








# canary



# canary





「cognito」

# services



## User Pools

Your own usernames,  
passwords, attributes, etc.



## Federated

Use SSO (Google, Facebook,  
Amazon, etc.)



## Sync

Synchronize data across  
devices

# user pools

## General settings

### Users and groups

Attributes

Policies

MFA and verifications

Advanced security <sup>beta</sup>

Message customizations

Tags

Devices

App clients

Triggers

Analytics

## App integration

App client settings

Domain name

UI customization

Resource servers

## Federation

Identity providers

Attribute mapping

# user information



## users and groups

List and manage users and/or groups and see details



## attributes

Define what a user looks like

# security



## policies

Password length, special characters, etc.



## mfa and verifications

MFA  
phone and email verification



## advanced

Intelligent security features



# customization



## message

Customize messages to users



## triggers

Lambda functions for user lifecycle



## devices

Remember devices

# other



## app clients

Define clients of the user pool



## tags

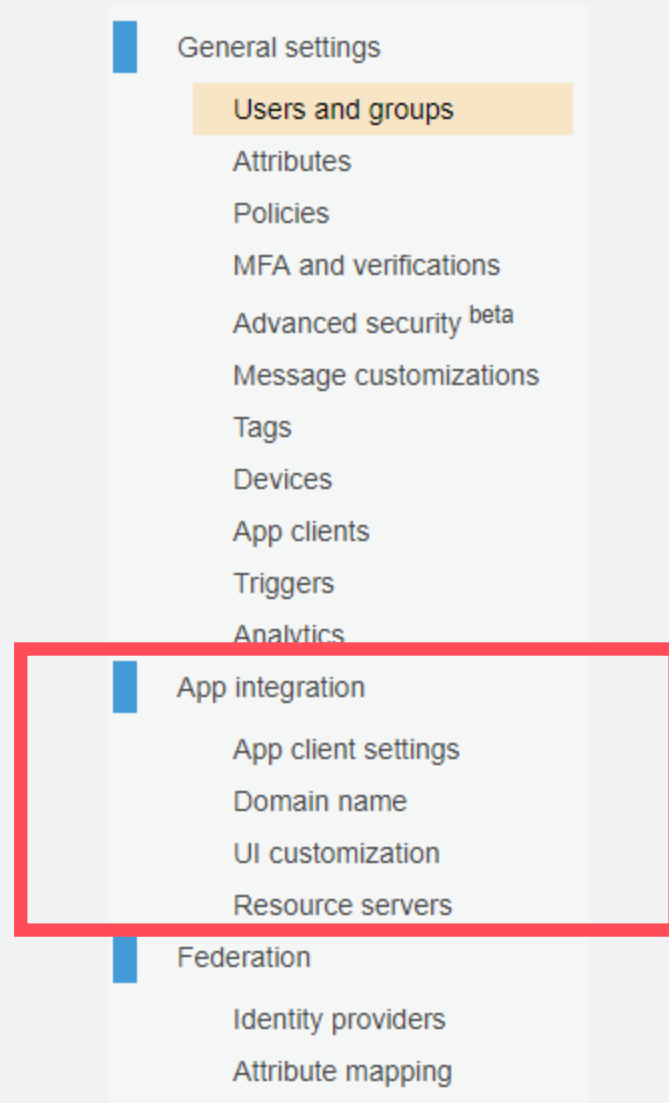
AWS tagging



## analytics

What are your users doing?


# OAuth 2.0



# federation?




# federation! and sync!





「cloudfront」

# what is a cdn?





# origin vs. behavior



## origin

Where the content comes from


S3, Web Server, API




## behavior

How to serve and cache a set of URIs

# caching



# invalidation



# other stuff

compression

Node.js at the edge